

Project Proposal Submission Form

Course: Distributed Systems

Instructor: Palak Mejpara

Semester: 25 Fall

Group ID / Team Name: Group 7

Submission Date: Oct / 28th / 2025

0. GitHub Repository

- **GitHub URL of project repository:** <https://github.com/Yuvi-5/Distributed-Weather-App>
(e.g., <https://github.com/your-team/yourproject>)
 - Ensure the repository is **public** (or shared with instructor) and will be used for all code, documentation, issues, and sprint tracking.
-

1. Team Information

Member Name	Student ID	Role (including Scrum Master)	Feature Owned
1. Dhruv Kalra	5143420	Backend Developer	Data Stream Handler
2. Janaranjanee Sendhil Kumar		Backend Developer	Alert & Broadcast Module
3. Dipal Nirmal		Frontend Developer	Visualization Dashboard
4. Minhajuddin Muhammad	5144183	Backend Developer	Aggregator & Forecast Engine
5. Yuvraj Singh Palh	5143317	Scrum Master	Edge Node Data Service

Note: Number of group members = number of major features

2. Project Title & Short Description

- **Project Title:**

Distributed Weather Data Aggregator

- **Short Description (2-3 sentences):**

A distributed, event-driven system that collects and broadcasts live weather updates from multiple nodes and predicts weather trends for the next seven days. Each node gathers real-time readings (temperature, humidity, and wind) and streams data to a central aggregator. The aggregator processes the data, detects anomalies, and uses an external API for 7-day forecasting. The dashboard visualizes all readings, alerts, and predictions in real time.

3. Features Summary

List all major features you plan to implement (excluding login/signup). Each feature corresponds to one team member.

Feature #	Feature Name	Brief Description	Assigned Member
1	Edge Node Data Service	Simulates or retrieves real-time local weather data (temperature, humidity, wind speed) and sends it to the broker.	
2	Data Stream Handler	Ensures reliable message delivery, manages buffering, and handles retries when nodes or network fail.	
3	Aggregator & Forecast Engine	Collects data from all nodes, computes regional statistics, detects anomalies, and integrates 7-day forecast from OpenWeatherMap API.	
4	Alert & Broadcast Module	Generates regional weather alerts and broadcasts live weather data streams to all connected dashboards.	
5	Visualization Dashboard	Real-time UI showing live readings, trends, node health, alerts, and 7-day forecast charts.	

4. Architecture Overview

- **Proposed System Architecture:**
(E.g., microservices, event-driven, peer-to-peer, using message queue, etc.)
- **Attach Architecture Diagram:**
The diagram must show:
 - Components/services/nodes
 - How they communicate (APIs, queues, RPC, etc)
 - Data flows, storage layers (database(s), cache(s))
 - Any external systems or integrations
- **Reasoning / Justification for the Architecture:**
Explain why you chose this architecture in terms of distributed-systems concerns: scalability, fault-tolerance, consistency, partitioning, latency, etc.

Proposed System Architecture:

A microservices-based event-driven architecture using a publish–subscribe model.

Edge Nodes: Produce local weather data and publish to Kafka/RabbitMQ.

Aggregator Service: Consumes data, processes statistics, stores them in MongoDB, and integrates 7-day forecast data from the OpenWeatherMap API.

Alert & Broadcast Module: Detects anomalies and pushes updates to WebSocket endpoints.

Dashboard: Subscribes to the broadcast for live display of metrics and forecasts.

Reasoning / Justification:


Scalability: Each node runs independently and new ones can join seamlessly.

Fault Tolerance: Queue-based communication ensures no data loss during downtime.

Consistency: Eventual consistency maintained across nodes and central aggregator.

Low Latency: WebSockets and message queues ensure fast live updates.

Partition Tolerance: Nodes can temporarily operate offline and resync later



5. Agile Methodology & Team Roles

- **Scrum Master:** Yuvraj
 - **Other team roles (e.g., Developer, QA, DevOps):**
 - _____
 - _____ Mentioned above
 - **Sprint Structure:**
 - Sprint length: 1 weeks
 - Number of sprints: 3
 - Sprint deliverables (example): _____
Sprint 1 → Node setup + data stream
 - **Tools for tracking & collaboration:**
(e.g., Trello, Jira, GitHub Projects, Slack, etc)
 - _____
Sprint 2 → Aggregator + forecast integration
 - _____
Sprint 3 → Dashboard + alert broadcasting
 - GitHub Projects (Kanban), Slack (communication), Google Meet (stand-ups)
 - **Stand-up Cadence:**
(e.g., daily at _____ time, 2 times per week)
 - _____
Twice weekly
 - **If your group has only two members,** tick here: *We will communicate separately with instructor prior to starting*
-

6. Timeline & Milestones

Milestone	Description	Target Date
Proposal submission & approval	Submit this document and get instructor approval	<u> / </u> / <u>Oct 28</u>
Sprint 1 end		<u>Nov</u> / <u>7</u> / <u>2025</u>
Sprint 2 end		<u>Nov</u> / <u>14</u> / <u>2025</u>
Mid-semester demo/update		<u>Nov</u> / <u>21</u> / <u>2025</u>
Final project submission	Working distributed system + report + demo	<u>End of/Sem</u> / <u> </u>

7. Risk Assessment & Mitigation

- **Risks:**
 1. Forecast API limitations or downtime.
 2. Data loss due to network or broker failure.
 3. Latency increase during concurrent updates.
- **Mitigations / Contingency Plans:**
 1. Cache API results and simulate fallback data.
 2. Use persistent message queues with retry mechanisms.
 3. Implement throttling and batching for heavy data flow.

8. Dependencies & External Tools

- List any **external libraries / frameworks / services / APIs** that your project depends on:
Frameworks: FastAPI / Node.js, Kafka or RabbitMQ, MongoDB / InfluxDB, React.
External APIs: OpenWeatherMap (for live and forecast data).
- Are there any **hardware or infrastructure** dependencies? (e.g., servers, cloud, nodes, IoT)
Docker Compose for distributed deployment.
- Are there any **known constraints** (time zones, availability, data privacy, third-party limits)?
Free API limits, local machine resource allocation.

9. Evaluation Metrics & Success Criteria

- **How will you evaluate your system's success?** (e.g., throughput, latency, fault recovery, number of nodes supported, user operations per second, etc.)

End-to-end latency (node → dashboard)	Forecast accuracy (%)
Throughput (messages/sec)	Alert latency and reliability
- **What are your target metrics/goals?**

Latency < 300 ms	0% message loss with ≥ 5 active nodes
Forecast accuracy $\geq 85\%$	Fault recovery within 5 seconds
- **How will you demonstrate/measure these?**

Use system logs, real-time dashboard metrics, and stress testing under simulated failures.

10. Approval Section (Instructor)

- Instructor Comments:

- Approved: Yes No
 - Date: ____ / ____ / ____
 - Signature: _____
-

Instructions to students:

1. Fill in all sections above completely and accurately.
2. Provide the architecture diagram as required in Section 4.
3. Submit this form **before you begin implementation.**
4. Do **not** start coding, designing in detail or building the system until you receive approval.